

# Deploying an OpenStack infrastructure on a single server with Ansible

James Freeman

April 2022



# Who am I?

- Technical Account Manager at AWS
- Empowerment Coach
- Reiki Master Teacher
- Autism Dad
- Inventor
- GTD Enthusiast
- Published Author (Ansible)



<https://github.com/jamesfreeman959>



[linkedin.com/in/jamesfreeman959](https://www.linkedin.com/in/jamesfreeman959)



[twitter.com/jamesfreeman959](https://twitter.com/jamesfreeman959)

# A horse walks into a bar...

- ▶ Why the long title?
  - ▶ It didn't fit on the title slide!
- ▶ “Digital Twinning with Ansible and cloud-init to support destructive infrastructure testing”
- ▶ Let me tell you a story

# In a prior role, the CEO says to me...

- ▶ Build me OpenStack
- ▶ A simple enough request in itself(?)
- ▶ All-in-one nodes are easy and quick to build
  - ▶ Simple
  - ▶ Effective for learning the basics
  - ▶ Bear no relation to real life
- ▶ Oh by the way...
  - ▶ Here's one single blade to build it on
  - ▶ With mechanical storage

# Digital Twinning

- ▶ First used at NASA
  - ▶ Full scale mock-ups of space capsules
  - ▶ Paved the way for digital simulations
- ▶ Can be used in any context
- ▶ My mission:
  - ▶ To build an OpenStack environment with an architecture and network representative of something you might deploy in the real world



# ...with Ansible and cloud-init...

- ▶ Hypothesis:
  - ▶ I'm not going to get this right first time
  - ▶ Rebuilds are going to be a necessary corrective action
    - ▶ Think Docker-like behaviour
  - ▶ I don't want to waste hours performing basic tasks
    - ▶ Building/deploying blank operating system images
    - ▶ Setting IP addresses and initial passwords
    - ▶ Deploying software and configuration

# Why Ansible?



- ▶ Yes, existing familiarity, but also...
- ▶ Ansible overtook Chef and Puppet last year as top cloud configuration management tool<sup>1</sup>
- ▶ Agentless
  - ▶ Uses SSH to communicate with nodes
  - ▶ Works well with switches
- ▶ Idempotent
- ▶ Self-documenting code
- ▶ Vault for securing sensitive data
- ▶ See also - the openstack-ansible project

1. Source: <https://www.techrepublic.com/article/ansible-overtakes-chef-and-puppet-as-the-top-cloud-configuration-management-tool/>

# Why cloud-init?



- ▶ Ansible is great if:
  1. Your operating system image has an IP address and appropriate network configuration
  2. You can authenticate with your operating system image when it boots
  3. (Linux nodes): Python is installed
- ▶ Simple to achieve, yet tedious if performed at scale and repetitively
- ▶ The solution: Couple publicly available cloud Ubuntu images with cloud-init
  - ▶ cloud-init is baked in
  - ▶ Each node has its own ISO image
  - ▶ Configures bonding, IP addressing, routing, initial password, SSH keys
  - ▶ Even installs Python!



# I have but one blade to give...

- ▶ Fairly clear that I'm going to need a hypervisor
- ▶ Requirements:
  - ▶ Free
  - ▶ Fast
  - ▶ Flexible
  - ▶ Support nested virtualization (CPU support is present)
- ▶ Linux KVM seemed like the obvious choice

# What about the switch layer?

- ▶ Must run on Linux KVM
- ▶ Preferably modern, and most importantly, free!
- ▶ Cumulus Networks (now NVIDIA)
  - ▶ Open standards based
  - ▶ Linux!
  - ▶ Runs on over 130 hardware platforms
    - ▶ including libvirt!
  - ▶ Capable of building modern web-scale network architectures

# My brain is melting

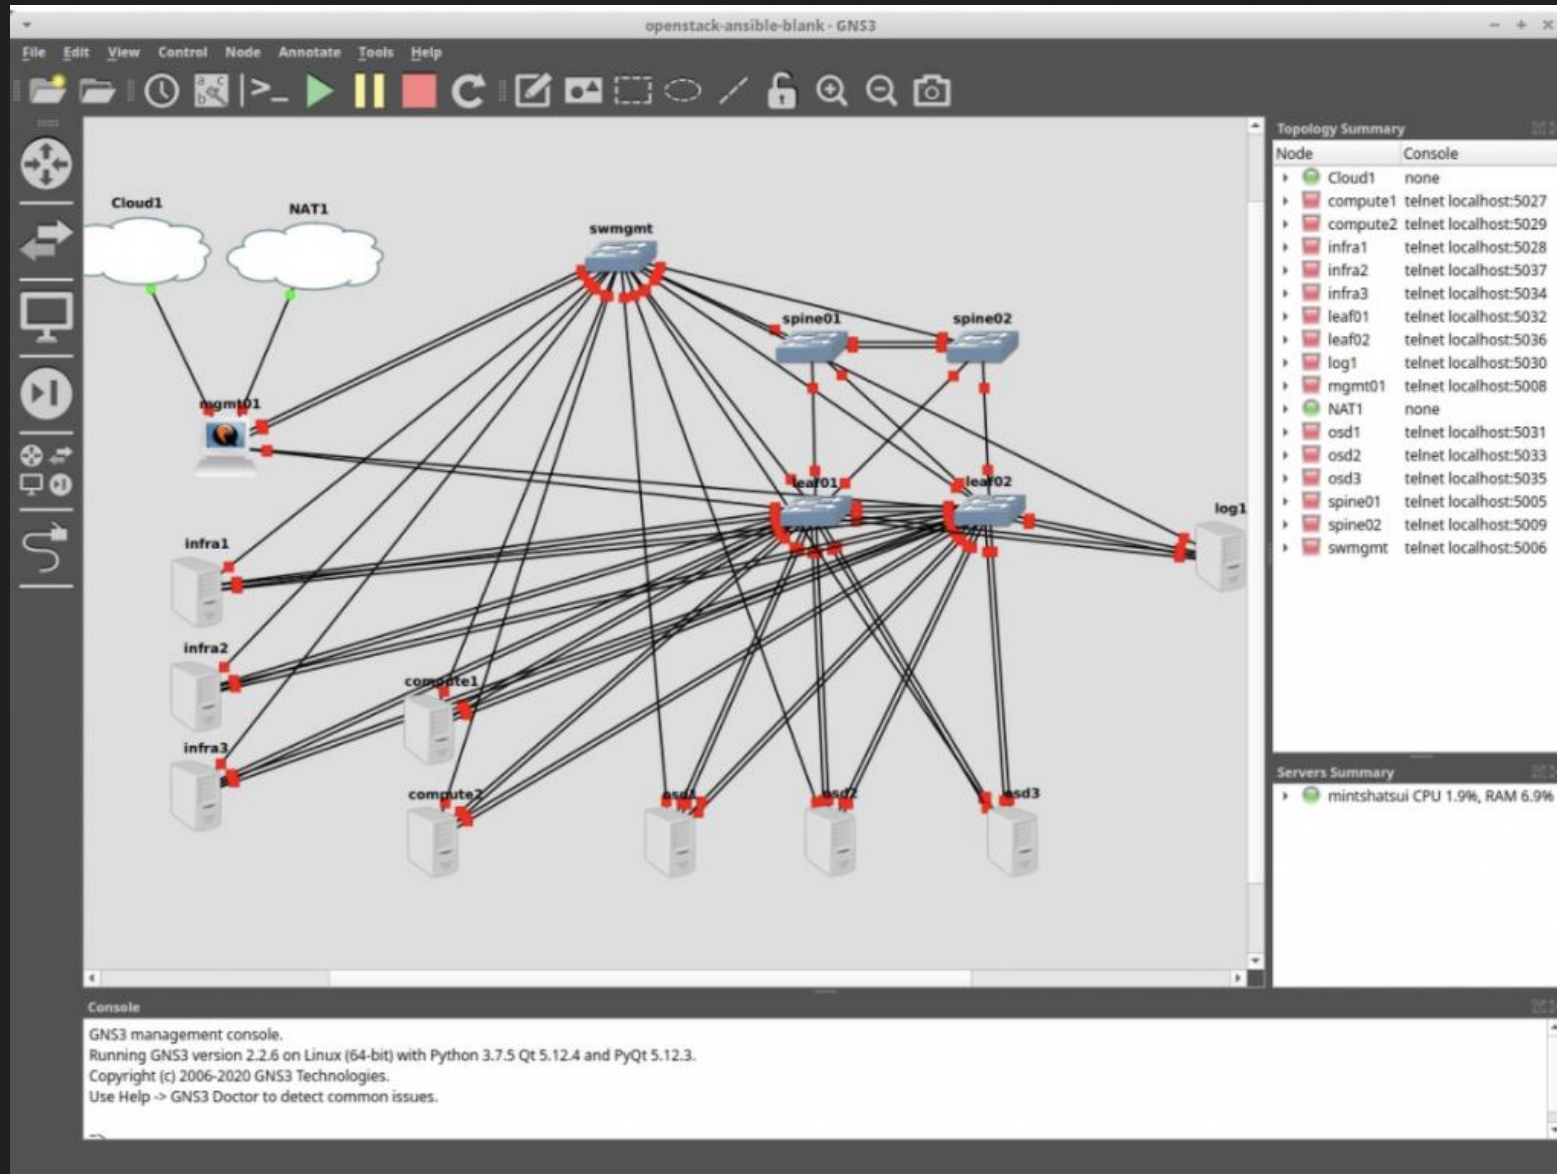
- ▶ Imagine a highly available deployment of OpenStack with Ceph for the backend storage
  - ▶ 10 Linux nodes, 5 switches
- ▶ Separate management, tunnel and storage networks
  - ▶ Using two different physical networks
  - ▶ VLAN tagging
  - ▶ VxLAN for the tunnel network
  - ▶ All ports bonded
  - ▶ Leaf and spine switch architecture for redundancy
- ▶ Now describe that in XML for libvirt

# Heresy in the form of GNS3



- ▶ I want a GUI to map out the network
- ▶ “I wonder if GNS3 would work...?”
- ▶ It turns out that GNS3:
  - ▶ Uses libvirt for to create and run virtual machines
  - ▶ Is designed to consume cloud images
  - ▶ Can directly support the use of ISO files for cloud-init
  - ▶ Supports extensive customization of VM's (e.g. adding nested virtualization)
  - ▶ Stores VM data in standard QCOW2 images that you can manipulate with qemu-img
- ▶ Oh, and it has a GUI

# The end result





# The end result

- ▶ It works! (after much tweaking)
- ▶ A representation of a possible real world OpenStack architecture
- ▶ All deployed with Ansible and cloud-init

# What does this actually mean?

- ▶ The playbooks and cloud-init data used to build the environment could be re-used, as is, to produce a real-world OpenStack infrastructure just like this
  - ▶ A real digital twin!
- ▶ The GNS3 network is isolated, so you could test with real world IP addresses and MAC addresses
- ▶ This technique can be applied to building just about any technology
  - ▶ Very little runs on just one server in production any more
- ▶ Perfect as a training tool!
  - ▶ Combines visual appeal with real-world like software deployments

# What else can we do...?

- Let's break something...



# Destructive testing ideas

- ▶ Network link failure
- ▶ Switch failure
- ▶ Node failure
- ▶ Denial of service attack
- ▶ Brute force attack
- ▶ Data corruption
- ▶ Backup and recovery techniques
- ▶ Business Continuity Planning

# Yet...

- ▶ You can back up and restore your entire digital twin with just a few commands or clicks
- ▶ GNS3 even supports snapshots of the entire environment!
- ▶ You get to blow it up more than once!



# A worked example

- ▶ In this test setup, we have 3 OSD nodes
  - ▶ All with a dedicated data disk
- ▶ Let's pull the plug on a node - does the cluster keep running as expected?
- ▶ Ok - now whilst that node is shut down, let's corrupt the disk
  - ▶ Idea - let's use dd from copy from /dev/random to the QCOW2 image backing the data
- ▶ Now power up the OSD node - what happens?
- ▶ Can we recover the node according to the standard recovery procedures?

# Limitations

- ▶ This is a big step forwards
- ▶ Moves us away from test environments not being representative of production
- ▶ But...
  - ▶ It's not fast compared to real hardware
    - ▶ Especially nested virtualization
  - ▶ No hardware specific testing
    - ▶ e.g. RAID controller, low level Ethernet setup, FibreChannel
  - ▶ You won't do this on your laptop

# Things I learned from this experience

- ▶ The Linux community is amazing, and those staffing the OpenStack Ansible IRC channel are so helpful
- ▶ Do throw hardware at this
  - ▶ I purchased a second hand workstation to run this on
  - ▶ 2 x Xeon (16c/32t), 128GB ECC RAM, 1TB SSD
- ▶ Do use virtio for disks and network cards wherever possible
- ▶ Don't attempt this setup with GNS3 on anything other than Linux
  - ▶ Windows = nested virtualization inside a VM

# Moving Forwards

- ▶ I now have a whole suite of lab environments built in this manner
- ▶ Configurations have evolved and become more streamlined
- ▶ Client demos are performed using this
  - ▶ People like the visual nature of the architecture
- ▶ I use this to train myself on new technologies
- ▶ Perhaps we can all learn more using these techniques?

Questions?





# Thank you!

- ▶ For more information, please reach out to me:



<https://github.com/jamesfreeman959>



[linkedin.com/in/jamesfreeman959](https://www.linkedin.com/in/jamesfreeman959)



[twitter.com/jamesfreeman959](https://twitter.com/jamesfreeman959)

- ▶ Code is available here:
  - ▶ <https://github.com/jamesfreeman959/gns3-cumulus-openstack-ansible>